

A High Efficiency MAC Protocol for WLANs: Providing Fairness in Dense Scenarios

Luis Sanabria-Russo, Jaume Barcelo, Boris Bellalta, Francesco Gringoli

Abstract—Collisions are a main cause of throughput degradation in WLANs. The current contention mechanism used in IEEE 802.11 networks is called Carrier Sense Multiple Access with Collision Avoidance (CSMA/CA). It uses a Binary Exponential Backoff (BEB) technique to randomise each contender attempt of transmitting, effectively reducing the collision probability. Nevertheless, CSMA/CA relies on a random backoff that while effective and fully distributed, in principle is unable to completely eliminate collisions, therefore degrading the network throughput as more contenders attempt to share the channel. Carrier Sense Multiple Access with Enhanced Collision Avoidance (CSMA/ECA) is able to create a collision-free schedule in a fully distributed manner using a deterministic backoff after successful transmissions. Hysteresis and Fair Share are two extensions of CSMA/ECA to support a large number of contenders in a collision-free schedule. CSMA/ECA offers better throughput than CSMA/CA and short-term throughput fairness.

This work describes CSMA/ECA and its extensions. Additionally, it provides the first evaluation results of CSMA/ECA with non-saturated traffic, channel errors, and its performance when coexisting with CSMA/CA nodes. Furthermore, it describes the effects of imperfect clocks over CSMA/ECA and present a mechanism to leverage the impact of channel errors over collision-free schedules. Finally, experimental results on throughput and lost packets from a CSMA/ECA implementation using commercial hardware and open-source firmware are presented.

Index Terms—CSMA/ECA, WLAN, MAC, Collision-free, Implementation.

I. INTRODUCTION

Wireless Local Area Networks (WLANs or WiFi networks [1]) are a popular solution for wireless connectivity, whether in public places, work environments or at home. This technology works over an unlicensed spectrum in the Industrial, Scientific and Medical (ISM) radio bands (at around 2.4 or 5 GHz), offering a good tradeoff between performance and costs, which is a main reason for its popularity.

The Medium Access Control (MAC) scheme used in WLANs is based on Carrier Sense Multiple Access with Collision Avoidance (CSMA/CA) protocol. It has been widely adopted by manufacturers and consumers, making it inexpensive to implement and an ubiquitous technology. Nevertheless, the ever-growing throughput demands from upper layers are faced with a bottleneck at the WLANs' MAC [2], which

by its nature is prone to collisions that degrade the overall performance as more nodes join the network.

The research community has pushed forward many alternatives to the current MAC in WLANs [3]–[13], but when a proposal deviates too much from CSMA/CA, or some time-critical operations are modified, its hardware implementation as part of WLANs' MAC often becomes unlikely [14], with the standardisation process taking many years without certainty of approval [2].

A CSMA/CA replacement should be able to provide advantages in terms of throughput, spectrum efficiency and number of supported contenders. All of the aforementioned without sacrificing short-term throughput fairness. Furthermore, it must also serve existing users, which means it has to be backwards compatible.

A suitable candidate, and the one to be evaluated in this work, is called Carrier Sense Multiple Access with Enhanced Collision Avoidance (CSMA/ECA) [7]. It is capable of attaining higher throughput than CSMA/CA by making a simple modification to the contention mechanism. In CSMA/ECA, nodes use a deterministic backoff after successful transmissions, constructing a collision-free schedule among successful contenders in a fully decentralised way. This backoff mechanism ensures that more channel time is spent on successful transmissions rather than recovering from collisions, thus increasing the throughput of the network. Further enhancements (or extensions), like *Hysteresis* and *Fair Share* [15] allow CSMA/ECA to support many more contenders in a collision-free schedule. Moreover, CSMA/ECA and its extensions are designed for allowing nodes to transmit as frequently as possible while keeping an even distribution of the available bandwidth among users.

The 802.11 High Efficiency WLAN (HEW) Study Group (SG) envisions very crowded scenarios as one of the future challenges for WLAN protocols [16], specifically those usually encountered at stadiums or conference rooms. Further, if the need for serving many users is combined with the increased throughput demands from the upper layers, a performance improvement at the MAC becomes paramount. Although many studies have been made analysing the performance of CSMA/ECA [7], [8], [15], [17], there are still several open aspects that require further attention and additional insight to consider CSMA/ECA as a potential CSMA/CA replacement. Namely, neither assesses the protocol's backwards compatibility property or its behavior under non-saturated traffic conditions while serving many users. Furthermore, the impact of channel errors and imperfect clocks over the deterministic backoff mechanism is also lacking.

This work was partially supported by the Spanish government, through the project CISNETS (TEC2012-32354).

L. Sanabria-Russo (luis.sanabria@upf.edu), J. Barcelo (jaume.barcelo@upf.edu) and B. Bellalta (boris.bellalta@upf.edu) are with the Department of Information and Communication Technologies, Universitat Pompeu Fabra, Barcelona, Spain.

F. Gringoli (francesco.gringoli@unibs.it) is with the Dipartimento di Elettronica per l'Automazione, Università degli Studi di Brescia, Brescia, Italy.

This paper fills those gaps by extending [15], and consolidates the push for CSMA/ECA as a potential replacement of CSMA/CA for next generation WLANs. In detail, this paper provides the following contributions:

- First results on the achievable throughput and delay of CSMA/ECA with Hysteresis and Fair Share under non-saturated traffic conditions for very large number of nodes.
- The impact of imperfect clocks on CSMA/ECA with Hysteresis and Fair Share's deterministic backoff and its consequences on the achieved performance.
- **The effect of channel errors over collision-free schedules and how to solve the issue of converging rapidly to the biggest cycle length with the Schedule Reset mechanism.**
- Formulation of the throughput bounds for CSMA/ECA with Hysteresis and Fair Share.
- Coexistence and backwards compatibility with CSMA/CA nodes under different traffic conditions.
- First implementation of CSMA/ECA with Hysteresis in real hardware and the experimental results in a crowded WLAN testbed.

Results, derived from a modified version of the COST [18] simulator, show that CSMA/ECA with Hysteresis and Fair Share is capable of accommodating many users in collision-free schedules. Further, as tests in a mixed network with different proportions of CSMA/CA and CSMA/ECA nodes show, the aggregated throughput is higher than the observed in CSMA/CA-only networks without significant sacrifice from legacy stations.

Beyond simulations results, the implementation of CSMA/ECA prototypes [19]–[22] show that the construction of collision-free schedules using a deterministic backoff after successful transmissions is possible and results in a throughput increase by reducing the number of corrupted frames. **We then present the first real hardware implementation of CSMA/ECA with Hysteresis using the open firmware OpenFWWF [23]. Results show that CSMA/ECA is able of providing higher throughput and lower losses than CSMA/CA for the same number of contenders.**

An overview of similar distributed and collision-free MAC protocols for WLANs is provided in Section II. CSMA/ECA, as well as its extensions for allocating many contenders in a collision-free schedule are explained in Section III. Section IV details the simulation environment for testing CSMA/ECA, while Section V explains the results. An overview of CSMA/ECA real-hardware implementation and testbed description is compiled in Section VI, followed by a summary of the still missing CSMA/ECA features needed to become the next CSMA/CA replacement in Section VII. Conclusions are drawn in Section VIII.

II. RELATED WORK

Time in WLANs is divided into tiny empty slots of fixed length σ_e , collisions, and successful slots of length σ_c and σ_s , respectively. Collision and successful slots contain collisions or successful transmissions, making them several orders of magnitude larger than empty slots ($\sigma_e \ll \min(\sigma_s, \sigma_c)$). One

of the effects of collisions is the degradation of the network performance by wasting channel time on collisions slots.

Recent advances in the WLANs PHY [2], [24] push the research community towards the development of MAC protocols able to take advantage of a much faster PHY. By reducing the time spent in collisions nodes are able to transmit more often, which in turn translates to an increase in the network throughput. Further, the upcoming MAC protocols for WLANs should work without message exchange between contenders, that is, work in a fully distributed fashion in order to avoid injecting extra control traffic that may reduce the data throughput.

The following are MAC protocols for WLANs, distributed and capable of attaining greater throughput than CSMA/CA by constructing a collision-free schedule.

A. Zero Collision MAC

Zero Collision MAC (ZC-MAC) [25] achieves a zero collision schedule for WLANs in a fully distributed way. It does so by allowing contenders to reserve one empty slot from a predefined virtual schedule of N -slots in length. Backlogged stations pick a slot in the virtual cycle to attempt transmission. If two or more stations picked the same slot in the cycle, their transmissions will eventually collide. This forces the involved contenders to randomly and uniformly select other empty slot from those detected empty in the previous cycle plus the slot where they collided. When all M stations reserve a different slot, a collision-free schedule is achieved.

ZC-MAC is able to outperform CSMA/CA under different scenarios. Nevertheless, given that the length of ZC MAC's virtual cycle has to be predefined without actual knowledge of the real number of contenders in the deployment, the protocol is unable to provide a collision-free schedule when $M > N$. Furthermore, if N is overestimated ($N \gg M$), the fixed-width empty slots between each contender's successful transmission are no longer negligible and contribute to the degradation of the network performance¹. Additionally, Z-MAC nodes require common knowledge of where the virtual schedule starts/ends. This is not considered in CSMA/CA and constitutes an obstacle towards standardisation.

B. Learning-MAC

Learning-MAC [26] is another MAC protocol able to build a collision-free schedule for many contenders. It does so defining a *learning strength* parameter, $\beta \in (0, 1)$. Each contender starts by picking a slot for transmission s of the schedule n of length C at random with uniform probability. After a contender picks slot $s(n)$, its selection in the next schedule, $s(n+1)$, will be conditioned by the result of the current attempt. (1) and (2) extracted from [26] show the probability of selecting the same slot $s(n)$ in cycle $n+1$.

$$\left. \begin{aligned} p_{s(n)}(n+1) &= 1, \\ p_j(n+1) &= 0, \end{aligned} \right\} \quad \text{Success} \quad (1)$$

¹With a very large N and $M \leq N$, the overall aggregated throughput will start to decrease as $M \rightarrow N$. This is due to the very big period of time between successful transmissions. For the opposite case, that is $M > N$, the same effect is produced albeit with an increased number of collisions.

$$\left. \begin{aligned} p_{s(n)}(n+1) &= \beta p_{s(n)}(n), \\ p_j(n+1) &= \beta p_j(n) + \frac{1-\beta}{C-1}, \end{aligned} \right\} \quad \text{Collision} \quad (2)$$

for all $j \neq s(n)$, $j \in \{1, \dots, C\}$. That is, if a station successfully transmitted in $s(n)$, it will pick the same slot on the next schedule with probability one. Otherwise, it follows (2).

The selection of β implies a compromise between fairness and convergence speed, which the authors determined $\beta = 0.95$ to provide satisfactory results.

L-MAC is able to achieve higher throughput than the current MAC with a very fast convergence speed. Nevertheless, the choice of β suppose a previous knowledge of the number of empty slots ($C - N$, where N is the number of contenders), which is not easily available to the current MAC or may require a centralised entity [13].

Further extensions to L-MAC introduced an *Adaptive* schedule length in order to increase the number of supported contenders in a collision-free schedule. This adaptive schedule length is doubled or halved depending on the presence of collisions or many empty slots per schedule, respectively. The effects of reducing the schedule length may provoke a re-convergence phase which can result in short-term fairness issues. Moreover, L-MAC nodes also require common knowledge of the start/end of the schedule.

III. CARRIER SENSE MULTIPLE ACCESS WITH ENHANCED COLLISION AVOIDANCE (CSMA/ECA)

CSMA/ECA [7] is a fully distributed and collision-free MAC for WLANs. It differs from CSMA/CA in that it uses a deterministic backoff, $B_d = CW_{\min}/2$ after successful transmissions, where CW_{\min} is the minimum contention window of typical value $CW_{\min} = 16$. By doing so, contenders that successfully transmitted on schedule n , will transmit without colliding with other successful nodes in future cycles.

Collisions are handled as in CSMA/CA, which is described in Algorithm 1. In Algorithm 1, the node starts by setting the retransmissions counter and backoff stage to zero ($r \in [0, R]$ and $k \in [0, m]$ respectively, m is the maximum backoff stage of typical value $m = 5$, and R is the retransmissions limit with typical value $R = m + 1$), then generates a random backoff, B . When the Acknowledgement (*ack*) for a sent packet is not received by the sender a collision is assumed. Upon collision, the involved nodes will double their contention window by incrementing their backoff stage in one and use a random backoff, $B \in [0, 2^k CW_{\min}]$. This procedure is described between Line 10 and 13 of Algorithm 1.

Algorithm 2 provides an explanation of CSMA/ECA's deterministic backoff mechanism, which main difference with CSMA/CA (and therefore with Algorithm 1) relies on the selection of a deterministic backoff after a successful transmission (compare Line 18 in Algorithm 1 with Line 18 in Algorithm 2). Figure 1 shows an example of CSMA/ECA dynamics with four contenders.

In Figure 1, the *STA-#* labels represent stations willing to transmit. The horizontal lines represent a time axis with each number indicating the amount of empty slots left for

```

1 while the device is on do
2    $r \leftarrow 0$ ;  $k \leftarrow 0$ ;
3    $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
4   while there is a packet to transmit do
5     repeat
6       while  $B > 0$  do
7         wait 1 slot;
8          $B \leftarrow B - 1$ ;
9       Attempt transmission of 1 packet;
10      if collision then
11         $r \leftarrow r + 1$ ;
12         $k \leftarrow \min(k + 1, m)$ ;
13         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
14      until ( $r = R$ ) or (success);
15       $r \leftarrow 0$ ;
16       $k \leftarrow 0$ ;
17      if success then
18         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
19      else
20        Discard packet;
21         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
22    Wait until there is a packet to transmit;
```

Algorithm 1: CSMA/CA. r indicates the number of retransmission attempts, while R is the maximum retransmission attempts limit. When it is reached, the packet waiting for transmission is dropped.

the backoff to expire. Stations willing to transmit begin the contention for the channel by waiting a random backoff, B . The first outline highlights the fact that stations STA-3 and STA-4 will eventually collide because they have selected the same B . After recomputing the random backoff, STA-4's attempt results in a successful transmission, which instructs the node to use a deterministic backoff, $B_d = 7$ in this case. By doing so, all successful STAs will not collide among each other in future cycles.

Collision slots being orders of magnitude larger than empty slots degrade the network performance. When CSMA/ECA builds the collision-free schedule all contenders are able to successfully transmit more often, increasing the aggregated throughput beyond CSMA/CA's. Figure 2, shows the average aggregated throughput of CSMA/ECA, CSMA/CA and two of the collision-free protocols presented in Section II, namely L-ZC and L-MAC [26].

Referring to Figure 2, CSMA/ECA is able to achieve an aggregated throughput that goes beyond CSMA/CA up until the number of contenders (N) is greater than B_d ($B_d = 7$ in the case of the figure). Beyond this point, the network will have a mixed behavior relating to backoff mechanisms: some nodes will successfully transmit and use a deterministic backoff while others will collide due to the lack of empty slots and return to a random backoff. As more contenders join the network, CSMA/ECA performance will approximate to CSMA/CA's. This effect is also observed in L-ZC and L-

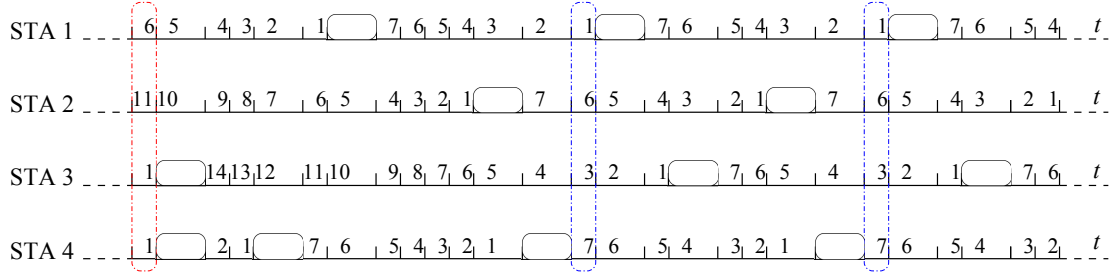


Fig. 1: An example of the temporal evolution of CSMA/ECA in saturation

```

1 while the device is on do
2    $r \leftarrow 0$ ;  $k \leftarrow 0$ ;
3    $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
4   while there is a packet to transmit do
5     repeat
6       while  $B > 0$  do
7         wait 1 slot;
8          $B \leftarrow B - 1$ ;
9       Attempt transmission of 1 packet;
10      if collision then
11         $r \leftarrow r + 1$ ;
12         $k \leftarrow \min(k + 1, m)$ ;
13         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
14      until ( $r = R$ ) or (success);
15       $r \leftarrow 0$ ;
16       $k \leftarrow 0$ ;
17      if success then
18         $B_d \leftarrow (2^k CW_{\min})/2 - 1$ ;
19         $B \leftarrow B_d$ ;
20      else
21        Discard packet;
22         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
23    Wait until there is a packet to transmit;

```

Algorithm 2: CSMA/ECA.

MAC, albeit at a larger number of nodes. This is because the virtual cycle used by these protocols (N for L-ZC and C for L-MAC in Section II) is greater than B_d . Specifically, the virtual cycle is of size equal to 16 (changing to $B_d = 16$ would only delay CSMA/ECA's throughput degradation to $N = B_d$).

The JFI curves in Figure 2 show the Jain's Fairness Index for all tested protocols. Showing a JFI equal to one suggests that the available throughput is shared evenly among all stations.

A. Supporting many more contenders

As was mentioned before, CSMA/ECA is only able to build a collision-free schedule if the number of contenders N , is less or equal than B_d . When $N > B_d$, collisions reappear.

To be able to attain a collision-free schedule even when the number of contenders exceeds B_d , we introduce *Hysteresis*. Hysteresis is a property of the protocol that instructs nodes not




Fig. 2: CSMA/ECA example in saturation: throughput (simulation parameters can be found in Section IV)

to reset their backoff stage (k) after successful transmissions, but to use a deterministic backoff $B_d = CW(k)/2$, where $CW(k) = 2^k CW_{\min}$. This measure allows the adaptation of the schedule length, admitting many more contenders in a collision-free schedule. This idea of a schedule is significantly different from the virtual schedule required by the protocols described in Section II, that is, CSMA/ECA with Hysteresis does not require a previous knowledge of the number of contenders, the result of previous transmissions or the start/end of the schedule, easing its implementation in real hardware.

Hysteresis enables CSMA/ECA nodes to have different schedules (B_d), carrying the undesired effect of unevenly dividing the channel time among contenders (i.e., some nodes will have to wait more in order to attempt transmissions).

This unfairness issue is solved by instructing nodes at backoff stage k to transmit 2^k packets on each attempt, thus proportionally compensating those nodes at higher backoff stages. This additional extension to CSMA/ECA is called *Fair Share*. CSMA/ECA with Hysteresis and Fair Share will be referred to as CSMA/ECA_{Hys+FS} in order to distinguish it from what was described until this point.

The idea of allowing the transmission of more packets to stations that transmit less often was initially proposed by Fang et al. in [26]. It was later adapted to CSMA/ECA_{Hys} and named Fair Share in [15]. Figure 2 shows the JFI for CSMA/CA as well as for CSMA/ECA_{Hys+FS}.

In Figure 2, the only curve deviating from JFI = 1 is


```

1 while the device is on do
2    $r \leftarrow 0$  ;  $k \leftarrow 0$ ;
3    $b \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
4   while there is a packet to transmit do
5     repeat
6       while  $B > 0$  do
7         wait 1 slot;
8          $B \leftarrow B - 1$ ;
9       Attempt transmission of  $2^k$  packets;
10      if collision then
11         $r \leftarrow r + 1$ ;
12         $k \leftarrow \min(k + 1, m)$ ;
13         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
14      until ( $r = R$ ) or (success);
15       $r \leftarrow 0$ ;
16      if success then
17         $B_d \leftarrow (2^k CW_{\min})/2 - 1$ ;
18         $B \leftarrow B_d$ ;
19      else
20        Discard packet;
21         $k \leftarrow 0$ ;
22         $B \leftarrow \mathcal{U}[0, 2^k CW_{\min} - 1]$ ;
23    Wait until there is a packet to transmit;

```

Algorithm 3: CSMA/ECA_{Hys+FS}

CSMA/ECA with Hysteresis (CSMA/ECA_{Hys}), suggesting an uneven partition of the channel access time among contenders (which is fixed with Fair Share).

Algorithm 3 shows an implementation of CSMA/ECA_{Hys+FS}, while an example of CSMA/ECA_{Hys+FS} with four contenders is shown in Figure 3. In the figure the first outline indicates a collision between STA-3 and STA-4, which will provoke an increment on both station's backoff stage ($k \leftarrow k + 1$). Once STA-4's random backoff expires, CSMA/ECA_{Hys+FS} instructs the station to transmit 2^k packets, and then use a deterministic backoff, $B_d = CW(k)/2$. The same behavior is followed by STA 3.

With Hysteresis and Fair Share, CSMA/ECA_{Hys+FS} is able to achieve greater throughput than CSMA/CA and for many more contenders, as shown also in Figure 2. In the figure, the CSMA/ECA_{Hys+FS} curve shows a greater throughput because collisions are eliminated and Fair Share allows nodes to send 2^k packets upon each transmission. This throughput increase is the result of aggregation via Fair Share, which also carries the effect of raising the average time between successful transmissions (see Section V-A3), which may affect delay-sensitive traffic, like gaming or live video/voice/tv streaming. Channel errors or unsaturated traffic conditions will disrupt any collision-free schedule, generate collisions and push all stations's backoff stages to the maximum value very quickly (contributing to the delay increase). This issue is leveraged with Schedule Reset mechanisms, presented in Section III-F.




Fig. 4: Throughput comparison with CSMA/CA_{MaxAg}: even though at low number of contenders CSMA/CA_{MaxAg} achieves greater throughput than CSMA/ECA_{Hys+FS}, collisions eventually degrade the throughput below CSMA/ECA_{Hys+FS}'s when the number of contenders increases past $N = 10$

B. The effects of Aggregation

Fair Share is an A-MPDU aggregation mechanism [27] that coupled with the collision-free schedule built by CSMA/ECA_{Hys} is able to provide short-term fairness. However, it also improves the throughput since the aggregation process makes the packet transmission more efficient by reducing overheads. The downside of Fair Share is that it may increase the time between two consecutive transmissions from the same node, which may affect negatively delay-sensitive applications such as gaming or high definition real-time video. An example of the duration of a transmission, $T(l_k)$, is provided by (5) in the following Section III-C.

In scenarios where short-term fairness and the time between consecutive transmissions are not relevant, Fair Share can be replaced by Maximum Aggregation (MaxAg), which will significantly improve the system throughput. In Maximum Aggregation all nodes aggregate as many packets as possible at every transmission, i.e., they send 2^m packets in each attempt.

Figure 4 shows the aggregated throughput for CSMA/ECA_{Hys+FS}, CSMA/ECA_{Hys+MaxAg}, CSMA/CA using the Fair Share mechanism (CSMA/CA_{FS}), and CSMA/CA_{MaxAg}. JFIs are shown below. Although CSMA/CA_{FS} and CSMA/CA_{MaxAg} perform aggregation, collisions degrade the aggregated throughput as more contenders attempt transmission. On the other hand, CSMA/ECA_{Hys+FS} is able to build a collision-free schedule and takes advantage of the aggregation provided by Fair Share, which opposed to just using CSMA/ECA_{Hys+MaxAg}, it is fair.

To summarise the effects of using aggregation:

- It increases the aggregated throughput: because nodes are able to send multiple packet in each attempt, the system throughput is increased. Moreover, Fair Share compensates those nodes at higher backoff stages to ensure throughput fairness.




Fig. 5: Upper and Lower throughput bounds for CSMA/ECA_{Hys+FS}

D. Clock drift issue in decentralized collision-free MAC protocols

CSMA/ECA relies on stations being able to correctly count empty slots and consequently attempt transmissions in the appropriate slot according to the backoff timer. Failure to do so may be caused by clock imperfections inside the Wireless Network Interface Cards (WNIC), which is commonly referred to as *clock drift*. As pointed out in [29], clock drift is a common issue that degrades the throughput in distributed collision-free MAC protocols like the ones reviewed in Sect. II.

While miscounting empty slots have no significant effect on CSMA/CA's throughput [29], it has a direct impact on CSMA/ECA. In a collision-free schedule with saturated CSMA/ECA contenders, a station miscounting empty slots will *drift* to a possibly busy slot, collide and force a re-convergence (if possible) to a collision-free schedule (see Sect. V-A2).

E. Channel errors

Failed transmissions due to channel errors are handled as collisions by CSMA/ECA_{Hys+FS}. Therefore, collision-free schedules under this type of channel model are shorter. In order to accelerate the convergence to collision-free schedules after channel errors CSMA/ECA_{Hys+FS} instruct nodes to *stick* to the current deterministic backoff even after *stickiness* number of consecutive failed transmissions. Stickiness is not a new concept to CSMA/ECA [13], it allows for a faster convergence towards a collision-free schedule, especially when performing under heavy channel errors.

Failed transmissions due to channel errors also means that a few moments of operation under a noisy channel can increase the contenders's deterministic backoff to its maximum value. Such event carries the undesired effects of increasing the time between successful transmissions, and a reduction in the overall system throughput due to shorter collision-free periods.

F. Schedule Reset Mechanism




Fig. 6: Example of the Schedule Reset mechanism.

CSMA/ECA_{Hys+FS} instructs nodes to keep their current backoff stage after a successful transmission (resetting it to zero only if reaching the retransmission attempt limit, see line 21 in Algorithm 3). This is done in order to increase the cycle length and provide a collision-free schedule for many contenders, which is desirable in dense scenarios.

Nevertheless, having a big deterministic backoff increases the time between successful transmissions. Furthermore, if not operating in a scenario with many nodes the empty slots between transmissions are not longer negligible and degrade the overall throughput of the system. For instance, if nodes withdraw from the contention their previously used slots will now be empty. On the other hand, in scenarios with channel errors contenders rapidly end up with the largest deterministic backoff. Nodes should be aware of this issue and pursue opportunities to reduce their deterministic backoff, B_d without sacrificing too much in collisions.

The *Schedule Reset* mechanism for CSMA/ECA_{Hys+FS} consists on finding the smallest collision-free schedule (if any) between a contender's transmissions and then change the node's deterministic backoff to fit in that schedule. Take a contender with $B_d = 32$ as an example. By listening to the slots between its transmissions, it is possible to determine the availability of smaller (and possibly) collision-free schedules.

Figure 6 shows the slots between the transmissions of a contender with $B_d = 32$. Starting from the left, the current $B_d = 0$ means that the first slot will be filled with the contender's own transmission. Each following slot containing either a transmission or a collision is identified with the number one, while empty slots are marked with a zero. Notice that the highlighted empty slots appear every eight slots, suggesting that a schedule reduction from $B_d = 32$ to $B_d^* = 8$ is possible². The schedule change is performed after the contender's next successful transmission.

1) *Keeping collision-free schedules:* Schedule Reset is described in Algorithm 4. Each contender starts filling a bitmap b of size B_d with the information observed in each passing slot during $\gamma = B_d/C$ consecutive successful transmissions (sxTx in Algorithm 4); where $C = 2^m CW_{\min}/2$ is the biggest deterministic backoff, and γ is referred to as the *Schedule Reset threshold* from this point forward.

Each bit t , $t \in \{0, \dots, B_d - 1\}$ in the bitmap is the result of a bitwise OR operation between its current value and the state of the corresponding slot. This is shown in Line 8 in Algorithm 4, where $\sigma_i(t)$ is a function that evaluates to 0 if the slot corresponding to $b[t]$ is empty in iteration i , $i \in$

² $B_d = 2^0 CW_{\min}/2 = 8$ (see Table I). So the change is made to the backoff stage, k . In this case from 2, to $k = 0$. This means that any new schedule must also be a power of two.

```

1 if  $sxTx == 1$  then
2   Initialising;
3   if  $B == 0$  then
4      $b[B_d] = \{0\}$ ;
5      $t = 0$ ;
6 if  $(sxTx > 0) \&\& (sxTx \leq \gamma)$  then
7   Filling the bitmap;
8    $b[t] \parallel \sigma_i(t)$ ;
9    $t++$ ;
10  if  $t == B_d - 1$  then
11     $t = 0$ ;
12 if  $sxTx == \gamma$  then
13   Analysing;
14    $sxTx = 0$ ;
15   for  $(j = 0; j < k; j++)$  do
16      $y = 2^j CW_{\min}/2$ ;
17      $isItPossible = 0$ ;
18     for  $(m = y; m < B_d; m++ = y)$  do
19        $isItPossible += b[m]$ ;
20     if  $isItPossible == 0$  then
21        $change = 1$ ;
22        $newStage = m$ ;
23       break;
24 if  $change == 1$  then
25   Making the change;
26    $k = newStage$ ;
27    $change = 0$ ;

```

Algorithm 4: Schedule Reset Mechanism for CSMA/ECA_{Hys+FS}. Every consecutive successful transmission increases the variable $SxTx$ by one, while a collision resets it to zero. The algorithm is called a *reset* when all smaller deterministic backoff are tested. On the other hand, when j in line 15 is initialised to $j = k - 1$, it is called a schedule *halving*

$\{1, \dots, \gamma\}$; or 1 otherwise. After γ iterations, the bitmap is evaluated (line 12 in Algorithm 4).

Schedule Reset tests all possible deterministic backoffs that are smaller than the current B_d (lines 15-19 in Algorithm 4), starting with the smallest one. If the corresponding bits in the bitmap are registered as empty, the process is stopped and the change to the new deterministic backoff is made (lines 20-27 in Algorithm 4). Otherwise, the process is restarted after the next successful transmission.

Using Figure 6 as an example, we can see that the sum $\sum_t b[t] = 0$, for $t \in \{8, 16, 24\}$. This means that the transmission slots corresponding to a deterministic backoff, $B_d^* = 8$ are free, therefore a change of schedule is possible. In case of suffering a collision immediately after applying the schedule reduction, the node reverses the changes made by Schedule Reset before handling the collision.

2) *Facing channel errors*: the default γ ensures that the bitmap registers all the transmission slots in the network

(assuming saturated traffic and a perfect channel), providing enough information for performing the schedule reduction without disrupting collision-free schedules.

Nevertheless, this γ can be rendered too conservative when in the presence of channel errors. This condition produces Schedule Reset bitmaps that do not contain enough information for successfully avoiding collisions after the schedule reduction. This is also true for any $\gamma^* > \gamma$.

This effect is leveraged by setting $\gamma = 1$, meaning that the bitmap is filled with the information of slots between only two consecutive transmissions. This measure increases the frequency of schedule reset attempts.

G. Backwards compatibility and coexistence

CSMA/ECA_{Hys+FS} springs from a modification to CSMA/CA's backoff mechanism. It keeps the range of values CSMA/CA nodes use to draw a random backoff (i.e., use the same CW_{\min} and CW_{\max}), allowing CSMA/ECA_{Hys+FS} contenders to coexist with CSMA/CA nodes in the same network. Further, the selection of CSMA/ECA_{Hys+FS}'s deterministic backoff, B_d , is the expected value for the current backoff stage k ($B_d := \lceil E[0, CW(k)] \rceil$) [15], which ensures fairness among contenders. An overview of the attained throughput for different proportions of CSMA/ECA_{Hys+FS} and CSMA/CA nodes is presented in Section V-C.

IV. SIMULATION SCENARIO

This section provides the simulation parameters for testing CSMA/ECA_{Hys+FS} under two different traffic conditions, namely saturated and non-saturated. **We also provide details on how channel imperfections are modelled and what are its effects over the transmissions.** Further, the simulation of the clock drift effect, and the coexistence with CSMA/CA are also subjects to be addressed in this section.

A. Scenario details

Results are obtained by running multiple simulations over a modified version of the COST simulator [18], available at [30]. PHY and MAC parameters are detailed in Table I. Some assumptions were made in order to test the performance at the MAC layer:

- Unspecified parameters follow the IEEE 802.11n (2.4 GHz) standard.
- All nodes are in communication range.
- No Request-to-Send (RTS) or Clear-to-Send (CTS) messages are used.
- Collisions take as much channel time as successful transmissions.

The aforementioned assumptions ensure that the simulation results are just effects of the MAC behaviour. If not mentioned otherwise, results are derived from **20 simulations of 100 seconds in length, each one with a different seed.** Figures also show the standard deviation.

TABLE I: PHY and MAC parameters for the simulations
REVIEW

PHY	
Parameter	Value
PHY rate	65 Mbps
Empty slot	9 μs
DIFS	28 μs
SIFS	10 μs
MAC	
Parameter	Value
Maximum backoff stage (m)	5
Minium Contention Window (CW_{\min})	16
Maximum retransmission attempts	6
Data payload (Bytes)	1024
MAC queue size (Packets)	1000

B. Saturated and Non-saturated stations

A saturated station always has packets in its MAC queue. This is modelled by setting the packet arrival rate to the MAC queue (Δ_{PAR}) to a value greater than the achievable throughput. To ensure saturation, stations are set to fill their MAC queue at $\Delta_{\text{PAR}} = 65$ Mbps, which is purposefully greater than the effective capacity of the channel.

To evaluate the performance under non-saturated conditions, stations need to be able to empty their MAC queues. To do so, the packet arrival rate to the MAC queue is set to $\Delta_{\text{PAR}} = 1$ Mbps. These values of Δ_{PAR} have proven to produce the desired effects.

C. Performance under clock drift

Clock drift is simulated by setting a drift probability, p_{cd} . Each station has a probability of $p_{cd}/2$ of miscounting one slot more, and $p_{cd}/2$ of miscounting one slot less. This approach follows the one proposed by Gong et. al in [29].

D. Channel errors

Channel errors are modelled by assigning a probability of a packet being corrupted by the channel, $p_e > 0$. That is, in a single packet transmission there is probability p_e that the transmission will not be acknowledged. If the transmission is an A-MPDU (like in the case of CSMA/ECA_{Hys+FS}), p_e will affect each packet individually and independently. Therefore, an A-MPDU transmission will be considered a failure if all packets in the aggregation are affected by the channel error probability.

All results are shown with stickiness equal to one (see Section III-D). That is, after a collision, CSMA/ECA_{Hys+FS} nodes will use a random backoff. Nevertheless, in Section V-D curves described as *dynStick* temporarily increase the node's stickiness to two after a successful reduction of the schedule (using a random backoff after two consecutive collisions). This increase is done in order to converge faster to collision-free schedules when in operating with channel errors.

E. Coexistence with CSMA/CA

To test the performance of CSMA/CA and CSMA/ECA_{Hys+FS} stations in the same network, simulations

are set with a CSMA/CA node density of 1/4, 1/2 and 3/4 of the total.

F. Applying Schedule Reset

A set of results under saturated conditions and channel errors applying Schedule Reset (see Section III-F) are provided. Some of the results are generated with a $\gamma = 1$, or *aggressive* Schedule Reset (aggr. in the figures). These settings provide the highest throughput under the tested conditions, and also in the real hardware implementations such as the one shown in Section VI.

V. RESULTS

A. Saturated nodes

In CSMA/CA, a large number of saturated nodes will normally be related to a high collision probability. This effect is in part the result of resetting the backoff stage after a successful transmission and the generation of a new random backoff. However, this scenario provides an advantageous condition to CSMA/ECA_{Hys+FS} nodes. In saturation, CSMA/ECA_{Hys+FS} nodes build a collision-free schedule and stick to their deterministic backoff as long as they transmit successfully, effectively eliminating collisions.

This section aims at overiewing the throughput of CSMA/CA and CSMA/ECA_{Hys+FS} in saturation, as well as the collision probability, the average time between successful transmissions and the effect of clock drift over the throughput.

1) *Throughput*: CSMA/ECA_{Hys+FS} nodes are able to build a collision-free schedule, use the channel more efficiently, and experience a throughput increase as seen in Figure 7a. Hysteresis allows the allocation of more contenders by increasing the length of the collision-free schedule, while Fair Share ensures an even distribution of the available throughput. This is reflected by the average backoff stage, which value increases with the number of contenders. In contrast, CSMA/CA throughput keeps decreasing due to an augmented number of collisions as the number of nodes increases (see Figure 7b). Further, Figure 7c shows the fraction of collision slots for CSMA/ECA_{Hys+FS} and CSMA/CA as simulation time passes. In the figure, is appreciated how the fraction of collision slots keeps decreasing once CSMA/ECA_{Hys+FS} reaches the collision-free schedule.

2) *Effect of clock drift over the achieved throughput in saturation*: Figure 8 shows the network aggregated throughput with 16 saturated stations and an increasing clock drift probability.

In Figure 8, a station has a clock drift probability equal to p_{cd} . Each station has a probability of $p_{cd}/2$ of miscounting one slot more, and $p_{cd}/2$ of miscounting one slot less. Because CSMA/CA is based on a random backoff, miscounting slots has no significant effect on the throughput. For the CSMA/ECA curve, it is possible to appreciate a slight decrease of the throughput as p_{cd} increases, caused by collisions due to the drift.

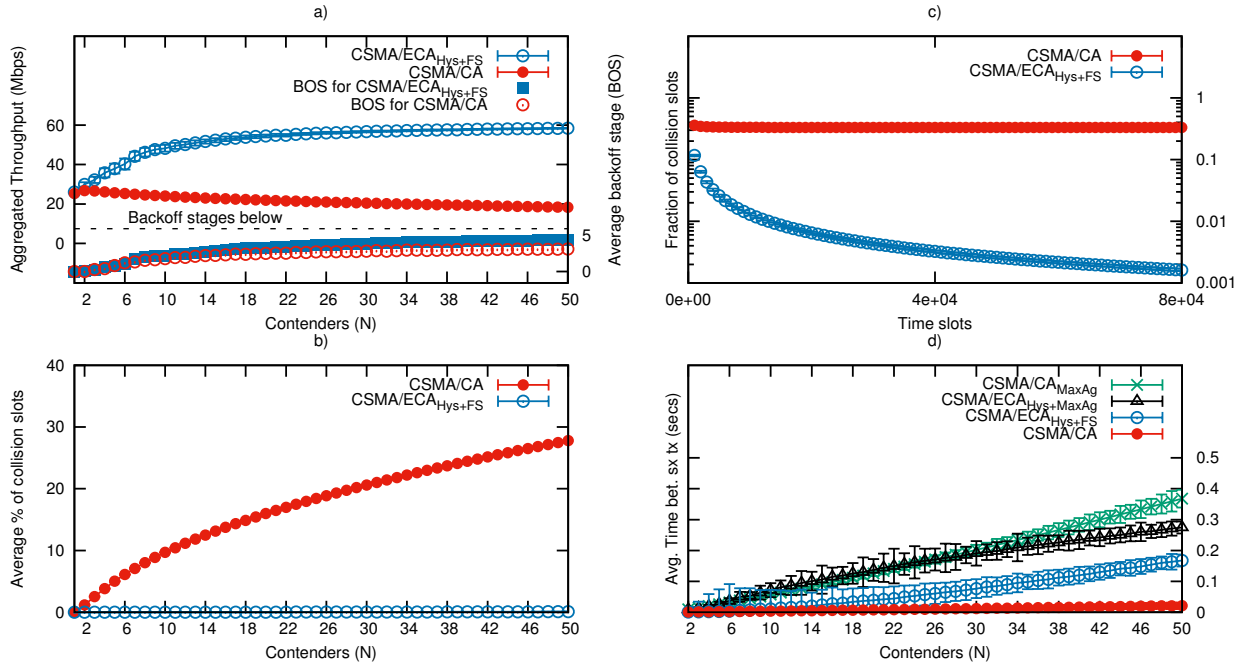


Fig. 7: Saturation traffic results: a) Throughput under saturated conditions; b) Average percentage of collision slots: fraction of time slots containing collisions; c) Evolution of the fraction of collision slots in a scenario with 70 saturated stations; d) Average time between successful transmissions (sx tx)




Fig. 8: Throughput when increasing the clock drift probability

The CSMA/ECA_{Hys+FS} curve in Figure 8 shows instead an increase of the aggregated throughput as p_{cd} grows. Collisions make CSMA/ECA_{Hys+FS} contenders to increment their backoff stage and aggregate packets for transmissions according to Fair Share, effectively increasing the throughput.

As it also can be appreciated in the figure, the average backoff stage for CSMA/ECA_{Hys+FS} contenders increases rapidly to its maximum value ($m = 5$), reducing the effect of clock drift over CSMA/ECA_{Hys+FS} nodes given that their transmissions would now be separated by more slots.

3) *Time between successful transmissions:* It is related to the elapsed time between the contention for transmission and

the reception of an ACK.

In Figure 7d all tests with maximum aggregation, namely CSMA/CA_{MaxAg} and CSMA/ECA_{Hys+MaxAg}, have an increased average time between successful transmissions. This is due to the multiple packets that are sent in each attempt. CSMA/CA_{MaxAg}, though, has an increased value due to collisions also taking longer channel time.

Although CSMA/ECA_{Hys+FS} has an increased average time between successful transmissions due to Fair Share, it has a lower metric when compared with the maximum aggregation curves in Figure 7d.

B. Non-saturated nodes

Emptying the MAC queue in CSMA/ECA_{Hys+FS} means that nodes will reset their backoff stage and use a random backoff when a new packet arrives at the queue, breaking the collision-free schedule (if any) for CSMA/ECA_{Hys+FS} contenders. The following show the impact over throughput, delay and time between successful transmissions when using CSMA/CA and CSMA/ECA_{Hys+FS} in non-saturated conditions.

1) *Throughput:* In Figure 9a, the aggregated throughput increases linearly for the CSMA/CA *non-saturated* curve until saturation is reached at around 22 nodes, where the throughput begins to degrade. The CSMA/ECA_{Hys+FS} *non-saturated* curve has a similar behavior, entering saturation at around 60 nodes. Further, at around 30 nodes we see an increase in the average backoff stage for CSMA/ECA_{Hys+FS} contenders which suggests an increment in collisions. This effect is shown in Figure 9b and Figure 9d, where at around 35 nodes

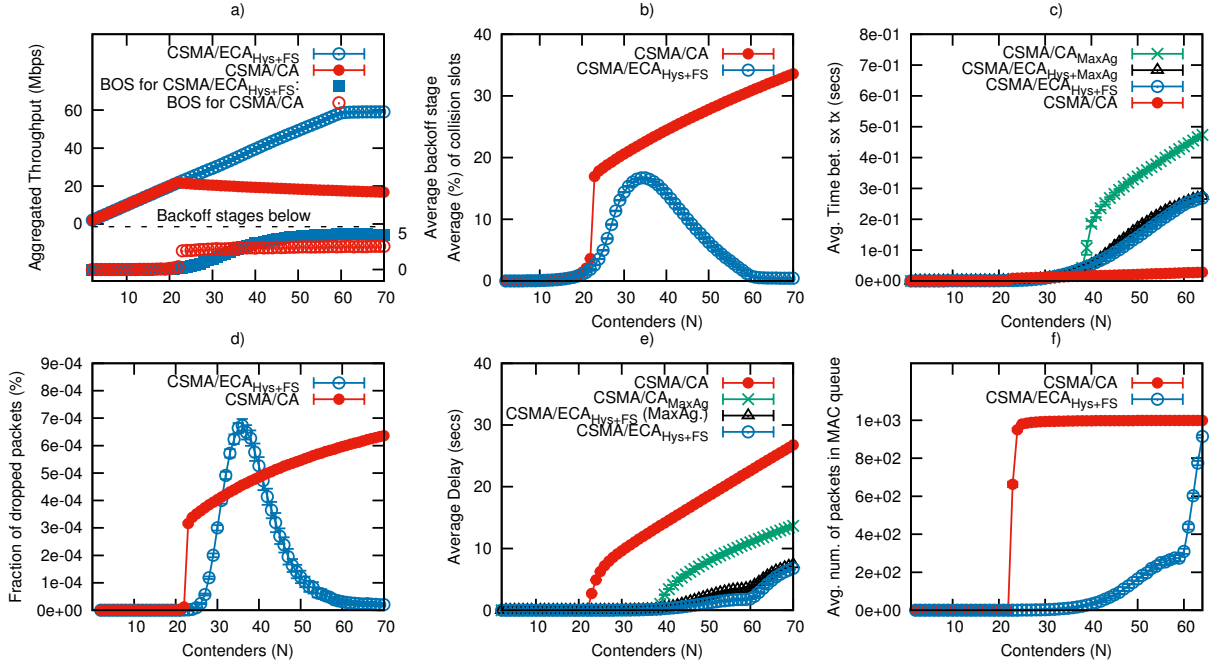


Fig. 9: Saturation traffic results: a) Throughput; b) Average percentage of collision slots: the fraction of time slots containing collisions; c) Average time between successful transmissions; d) Average fraction of dropped packets; e) Average system delay; f) Average number of packets in the MAC queue of a node

CSMA/ECA_{Hys+FS} contenders start colliding and dropping packets.

As indicated by Figure 9b, when $20 < N \leq 35$ CSMA/ECA_{Hys+FS} nodes suffer from an increasing number of collisions. This is due to nodes entering the contention with a random backoff every time a new packet arrives at their empty MAC queues. Beyond 35 contenders, the MAC queue of CSMA/ECA_{Hys+FS} nodes starts to fill up, as appreciated in Figure 9f, gradually allowing the construction of longer collision-free schedules due to CSMA/ECA_{Hys+FS} nodes getting saturated. This allows CSMA/ECA_{Hys+FS} to outperform CSMA/CA.

2) Delay: This metric refers to the elapsed time between the injection of a packet into the station's MAC queue and the reception of an ACK for such packet.

In Figure 9e, a rapid increase in the delay for CSMA/CA nodes is appreciated at the saturation point (around 20 contenders), whereas CSMA/ECA_{Hys+FS}'s delay is still low.

Further, with CSMA/ECA_{Hys+FS} the percentage of blocked packets from the MAC queue is lower than CSMA/CA or CSMA/CA_{MaxAg} (see Figure 10). This is due to the construction of a collision-free schedule which ensures that nodes transmit frequently (Hysteresis) and Fair Share which reduces the average delay given that more packets are transmitted and acknowledged with a single ACK.

As CSMA/ECA_{Hys+FS} nodes get saturated, the delay increases due to longer queueing and contention time (see the number of packets in the MAC queue for CSMA/ECA_{Hys+FS} nodes in Figure 9f and how it is related to the increase in delay shown in Figure 9e).




Fig. 10: Average fraction of blocked packets. When the MAC queue is full, packets coming from the upper layers must be discarded or blocked

Figure 9c shows the average time between successful transmissions. It is possible to see from the figure that when the saturation point is approached the average time between successful transmissions increases, resembling Figure 7d. Moreover, this effect is related to an increase in collisions: colliding nodes should begin a new contention for a retransmission, which may result in a successful transmission or another collision, the latter further increasing the average time between successful transmissions.




Fig. 11: Network throughput when composed by various proportions of CSMA/CA and CSMA/ECA_{Hys+FS} nodes




Fig. 13: Network throughput when composed by various proportions of non-saturated CSMA/CA and CSMA/ECA_{Hys+FS} nodes




Fig. 12: Average percentage of collision slots for the tested mixed network setups proportions




Fig. 14: Average percentage of collision slots for the tested mixed network setups proportions in non-saturated conditions

C. Coexistence with CSMA/CA

CSMA/ECA is thought to be an evolution of CSMA/CA given its similarities and the ability to coexist with the latter. This section provides simulation results for a setup of different proportions of CSMA/CA nodes in a network where there are also CSMA/ECA_{Hys+FS} contenders, that is: 1/4, 1/2 and 3/4 of the total nodes run CSMA/CA, while the rest uses CSMA/ECA_{Hys+FS}. This network configuration will be referred to as *mixed network setup* from here on.

1) *Throughput*: Figure 11 shows the network throughput for different proportions of CSMA/CA nodes in a mixed network setup.

In the figure it is appreciated how the mixed network setups curves lay between the CSMA/CA and CSMA/ECA_{Hys+FS} curves. As the proportion of CSMA/CA nodes decreases, the throughput increases as the result of a lower probability of collision, as can be seen in Figure 12. A similar behaviour is observed when testing the same proportion of nodes under non-saturated conditions. Figure 13 and Figure 14 show the average aggregated throughput and fraction of collisions slots

in a non-saturated mixed network setup.

As shown in Figure 11 and Figure 13, at a lower proportion of CSMA/CA nodes (1/4) the average aggregated throughput is the greatest among the tested mixed network setups. This is because collisions trigger Hysteresis and Fair Share in CSMA/ECA_{Hys+FS} nodes, lowering the number of times these nodes enter in a contention and reducing the overall collision probability when compared to an only CSMA/CA network (see Figure 12 and Figure 14). As the proportion of CSMA/CA nodes increases, the network throughput approximates to CSMA/CA.

Figure 15 shows the average throughput for CSMA/CA and CSMA/ECA_{Hys+FS} stations in a saturated mixed setup (half the contenders using CSMA/CA). It is possible to see that for a low total contenders ($N^* \leq 12$) CSMA/CA stations attain greater throughput than in a CSMA/CA-only network. This is because in the mixed network setup the other $N^*/2$ contenders with CSMA/ECA_{Hys+FS} use a deterministic backoff, leaving many empty slots between transmissions. This reduces the




Fig. 15: Average station throughput per MAC protocol in a saturated mixed network

chances of collisions for CSMA/CA nodes when compared with a CSMA/CA-only network and the same number of total contenders.

Still referring to Figure 15, for $N^* > 12$ periods of collision-free operation and the aggregation performed with Fair Share allows CSMA/ECA_{Hys+FS} to have larger channel time than CSMA/CA nodes, which throughput degrades even more than in CSMA/CA-only networks.

D. Saturation, channel errors and Schedule Reset

Figure 16 shows the average aggregated throughput of a non-saturated network with a channel error probability, $p_e = 0.1$ (see Section IV-D). We can see that the highest throughput, corresponding to CSMA/ECA_{Hys+FS} is lower than the one observed in a perfect channel (Figure 7a). Furthermore, curves with Schedule Reset (SR) seem to have lower aggregated throughput. This is because the backoff stage, and therefore the level of aggregation done by Fair Share, is constantly reduced.

Still assessing Figure 16 and focusing on the SR curves, we can see that both *SR aggr. halv. dynStick* curves (with and without Fair Share) have higher aggregated throughput than the rest. This Schedule Reset configuration shows higher aggregated throughput because it reduces the time between successful transmissions (see Figure 17) and maintains collision-free schedules for longer periods of time thanks to the increase in the node's stickiness.

Figure 17 shows the average time between successful transmissions for the same network setup. Even-though CSMA/ECA_{Hys+FS} still has a higher metric than CSMA/CA due to the aggregation done by Fair Share, Schedule Reset is able of reducing this metric by almost a 43%.

VI. REAL HARDWARE IMPLEMENTATION

We confirmed the simulations results experimentally with a testbed of real prototypes. This phase was mandatory to check whether the behaviour of the system is still correct in the presence of unexpected phenomena, like excessive channel noise,




Fig. 16: Average throughput for a saturated network with channel error ($p_e = 0.1$). Curves called "aggr" refer to values of $\gamma = 1$, while "halv" refers to a schedule halving (see Section III-F2 and Algorithm 4). *dynStick* refers to curves where the node's stickiness is temporarily increased to two after a successful reduction of the schedule




Fig. 17: Time between successful transmissions for a saturated network with channel error ($p_e = 0.1$). Curves called "aggr" refer to values of $\gamma = 1$, while "halv" refers to a schedule halving (see Section III-F2 and Algorithm 4). *dynStick* refers to curves where the node's stickiness is temporarily increased to two after a successful reduction of the schedule

or technical limitations of the hardware like imperfect node synchronisation; and to verify the implementation feasibility of the Schedule Reset mechanism with a real time horizon.

As CSMA/ECA_{Hys} works with the standard 802.11 PHY, we do not need the flexibility offered by complex development kits based on FPGAs like the WARP boards [31]: any architecture that allows a customisation of the channel access delay would fit our requirements. For this reason we chose a widely available 802.11b/g WNIC from Broadcom. Like many other cards, its behaviour is ruled by a firmware code that controls the underlying hardware (including the radio, carrier sense and FIFOs) in real time. In particular we used the BRCM4318KBFG

PCI card as it is supported by the OpenFWWF [23] open-source firmware, an alternative to the original code from the manufacturer that has been recently considered as development platform in several research projects [14], [22], [32]–[34].

Thanks to the extremely low price of this card (refurbished devices are widely available for less than \$10 each), it is possible to run experiments with very large testbeds. We used 25 PC-Engines Alix 2d2 nodes running Linux kernel 2.6.32 as stations and a more powerful computer with the same kernel as access point, all nodes being equipped with the aforementioned WNIC.

We built on the DCF protocol implemented in OpenFWWF and adapted the firmware to create collision-free schedules as described in Algorithm 2. The modification was straightforward: for every transmitted data frame the firmware sets an ACK-time-out alarm. Later, either at the expiration of the timer or when the acknowledgment frame is received, it executes a handler labelled `tx_contention_params_update`. It updates the contention window value `STATE_CW` and back-off counter according to the success/failure of the previous transmission attempt, just as in Algorithm 2. Implementing CSMA/ECA_{Hys} required just a modification specifying that a reset of the backoff stage (or `STATE_CW` in this case) is performed only when a packet is dropped or when the MAC queue empties.

To incorporate stickiness into the prototype we added a `STATE` to the system that can be either `STATE_DETERMINISTIC` or `STATE_RANDOM`, and a `STICKINESS` counter that we reset each time we have a success and decrement when failing. When the counter gets to zero we enter the random state. Upon a successful transmission we unconditionally enter the deterministic state and reset the stickiness counter to `DEFAULT_STICKINESS`.

For the Schedule Reset mechanism, we added a bitmap for monitoring the state of every single slot after a successful transmission. To index the current slot in the bitmap we used the hardware register `SPR_IFS_BKOFFDELAY`, which counts how many slots have still to come before the next transmission. The value of the register is decremented once per idle slot as requested by the DCF procedure, so the idea is to start with an empty bitmap and mark those slots for which the countdown procedure was interrupted.

To avoid spurious detection, instead of continuously checking the state of the channel, we rely on the execution of the `rx_plcp` handler, which is called each time a valid PLCP is detected. When this happens, we implicitly know that the backoff counter has been frozen at least $20\mu s$ ago, which is the time between the detection of the first short trailing sequence and the complete decoding of the PLCP signal data. In any case, the slot is marked as busy.

We have the option to keep filling the bitmap for up to `BITMAP_ROUND_BUILD` consecutive successful transmissions (γ in Section III-F) before checking if the central slot in the bitmap is available (a schedule *halving*), and if that is the case we perform a halving of the node's current schedule.

We performed four experiments for every number of contenders, starting from 1 up to 25. For every experiment each station establishes an iPerf [35] session and transmits saturated

TABLE II: PHY, MAC and other parameters for the real hardware experiments **REVIEW**

PHY	
Parameter	Value
PHY rate	48 Mbps
Empty slot	9 μs
DIFS	28 μs
SIFS	10 μs
IEEE 802.11g WiFi channel	14
MAC	
Parameter	Value
Maximum backoff stage (m)	5
Minium Contention Window (CW_{min})	16
Maximum retransmission attempts	6
Data payload (Bytes)	1470
Schedule Reset γ	1
Schedule Reset mode	halving, dynStick
TESTBED	
N	25
Distance between nodes and AP	8 m
Arrangement of nodes	Semicircle

UDP traffic towards a central AP, which also served as iPerf server for each flow. We used WiFi channel 14 in order to avoid interference from other networks. The aggregated throughput is derived from the iPerf logs, while the percentage of lost frames is reported by the firmware. This is known to be $(tx - sx)/tx$; where tx are the number of transmission attempts, and sx are the number of acknowledged frames.

Figure 18 shows the average aggregated throughput, while Figure 19 presents the average percentage of lost frames. Details of the testbed are shown in Table II.

CSMA/ECA_{Hys} has greater throughput due to its ability to avoid collisions more efficiently than CSMA/CA. Further, Schedule Reset (or halving in this case) prevents CSMA/ECA_{Hys} nodes from increasing too much the time between successful transmissions. Nevertheless, the implementation results reveal that there are other underlying factors that disrupt collision-free schedules. This is evidenced by the increasing percentage of lost frames followed by a throughput degradation in CSMA/ECA_{Hys}. The analysis of the factors that may disrupt collision-free schedules on real hardware implementations of CSMA/ECA_{Hys} is left as a future research topic.

VII. TRANSITIONING TOWARDS CSMA/ECA_{Hys+FS}

The current PHY enhancements considered by the HEW study group include higher modulation and coding schemes as well as full duplex radios, capable of receiving and transmitting at the same time, thus augmenting the achievable throughput. Furthermore, already existing technologies like Channel bonding and Multiple-User Multiple-Input Multiple-Output (MU-MIMO) also seek to increase the throughput. The latter by allowing the transmission of different packets to multiple destinations at the same time, while the former bonds multiple WiFi channels together in order to increase the available bandwidth. Using CSMA/ECA_{Hys+FS} alongside these features would provide enhanced performance by constructing collision-free schedules, thus substantially decreasing the time spent recovering from collisions.

Additionally, there are several features and scenarios still to




Fig. 18: Average aggregated throughput for a saturated network. Real hardware implementation results (see Table II)




Fig. 19: Average percentage of losses for a saturated network. Real hardware implementation results (see Table II)

be analysed for CSMA/ECA_{Hys+FS} networks. Part of what is left for future work is summarised in the following:

- The performance of a CSMA/ECA_{Hys+FS} network in presence of other neighbouring wireless networks and hidden nodes: transmissions from other network may negatively affect the maintenance of the collision-free schedule when not all devices in the CSMA/ECA_{Hys+FS} WLAN are able to listen to them. As a consequence, not all stations will pause the backoff accordingly, resulting in large slot drifts that in consequence will disrupt the collision-free schedule, approximating CSMA/ECA_{Hys+FS} performance to CSMA/CA's. A similar effect is expected when in the presence of hidden nodes.
- Traffic differentiation: although priorities using different contention windows in CSMA/ECA_{Hys+FS} proved to outperform CSMA/CA, other Enhanced Distributed Channel Access (EDCA) mechanisms like the Arbitration Inter-Frame Spacing (AIFS) would not work [36], whereas the following EDCA mechanisms would:
 - Transmission Opportunity (TXOP): stations with an

increased TXOP are able to transmit more packets. APs in CSMA/ECA_{Hys+FS} can use a big TXOP in order to transmit more than users.

- Multiple Queues: [36] provides performance metrics for two traffic categories. In order to transition to CSMA/ECA_{Hys+FS} a total of four access categories (AC) should be implemented, as in EDCA. Priorities to the multiple queues can be granted through different minimum and maximum contention windows.

CSMA/ECA_{Hys+FS} with multiple access categories is expected to provide better traffic differentiation in WLANs, mainly due to the elimination of collisions by using a deterministic backoff after each access category's successful transmission.

- Coexistence with EDCA stations in the same WLAN.

VIII. CONCLUSIONS

CSMA/ECA_{Hys+FS} is able to construct a collision-free schedule with many contenders. Taking advantage of this condition, the cumulative throughput experienced by CSMA/ECA_{Hys+FS} nodes goes beyond the achievable by CSMA/CA for any number of nodes. All of these while preserving throughput fairness. Further, as non-saturated CSMA/ECA_{Hys+FS} networks get crowded the cumulative throughput keeps increasing up to the saturation point, contrary to the throughput degradation seen in CSMA/CA networks.

As CSMA/ECA_{Hys+FS} is thought to be the next MAC for WLANs, coexistence with CSMA/CA nodes is paramount. Results show that coexistence is not an issue for any proportion of CSMA/ECA_{Hys+FS}/CSMA/CA users. Moreover, when the network is composed with a majority of CSMA/ECA_{Hys+FS} nodes, the cumulative throughput increases from CSMA/CA's.

To top it all, the real hardware implementation of CSMA/ECA_{Hys} with Schedule Reset clearly outperforms CSMA/CA. In itself, the prototype represents a strong indicator that the proposed enhanced collision avoidance could be considered as a viable replacement or evolution of the current MAC for WiFi.

REFERENCES

- [1] "IEEE Standard for Information Technology - Telecommunications and Information exchange between systems. Local and Metropolitan Area Networks - Specific requirements," *IEEE Std 802.11TM-2012*, p. 1646, 2012.
- [2] Perahia, E., "IEEE 802.11n Development: History, Process, and Technology," *Communications Magazine, IEEE*, vol. 46, no. 7, pp. 48–55, 2008.
- [3] V. Bharghavan, A. Demers, S. Shenker, and L. Zhang, "MACAW: a media access protocol for wireless LAN's," in *ACM SIGCOMM Computer Communication Review*, vol. 24, no. 4. ACM, 1994, pp. 212–225.
- [4] C. Wang, B. Li, and L. Li, "A new collision resolution mechanism to enhance the performance of IEEE 802.11 DCF," *IEEE Transactions on Vehicular Technology*, vol. 53, no. 4, pp. 1235–1246, 2004.
- [5] F. Cali, M. Conti, E. Gregori, and P. Aleph, "Dynamic Tuning of the IEEE 802.11 Protocol to Achieve a Theoretical Throughput Limit," vol. 8, no. 6, pp. 785–799, 2000.
- [6] A. Lopez-Toledo, T. Vercauteren, and X. Wang, "Adaptive Optimization of IEEE 802.11 DCF Based on Bayesian Estimation of the Number of Competing Terminals," vol. 5, no. 9, p. 1283, 2006.
- [7] J. Barcelo, B. Bellalta, C. Cano, and M. Oliver, "Learning-BEB: Avoiding Collisions in WLAN," in *Eunice*, 2008.

- [8] J. Barcelo, B. Bellalta, A. Sfairopoulou, C. Cano, and M. Oliver, "CSMA with Enhanced Collision Avoidance: a Performance Assessment," in *IEEE VTC Spring*, 2009.
- [9] Y. He, R. Yuan, J. Sun, and W. Gong, "Semi-Random Backoff: Towards resource reservation for channel access in wireless LANs," in *17th IEEE International Conference on Network Protocols*. IEEE, 2009, pp. 21–30.
- [10] J. Barcelo, A. Toledo, C. Cano, and M. Oliver, "Fairness and Convergence of CSMA with Enhanced Collision Avoidance (ECA)," in *2010 IEEE International Conference on Communications (ICC)*, may 2010, pp. 1–6.
- [11] M. Fang, D. Malone, K. Duffy, and D. Leith, "Decentralised learning MACs for collision-free access in WLANs," *Wireless Networks*, vol. 19, pp. 83–98, 2013. [Online]. Available: <http://dx.doi.org/10.1007/s11276-012-0452-1>
- [12] K. Hui, T. Li, D. Guo, and R. Berry, "Exploiting peer-to-peer state exchange for distributed medium access control," in *Information Theory Proceedings (ISIT), 2011 IEEE International Symposium on*. IEEE, 2011, pp. 2368–2372.
- [13] J. Barcelo, B. Bellalta, C. Cano, A. Sfairopoulou, and M. Oliver, "Towards a Collision-Free WLAN: Dynamic Parameter Adjustment in CSMA/E2CA," in *EURASIP Journal on Wireless Communications and Networking*, 2011.
- [14] I. Tinnirello, G. Bianchi, P. Gallo, D. Garlisi, F. Giuliano, and F. Gringoli, "Wireless MAC processors: Programming MAC protocols on commodity Hardware," in *INFOCOM, 2012 Proceedings IEEE*, march 2012, pp. 1269–1277.
- [15] L. Sanabria-Russo, A. Faridi, B. Bellalta, J. Barceló, and M. Oliver, "Future Evolution of CSMA Protocols for the IEEE 802.11 Standard," in *2nd IEEE Workshop on Telecommunications Standards: From Research to Standards*. Budapest, Hungary, 2013.
- [16] IEEE 802.11 Study Group. (2014) Status of IEEE 802.11 HEW Study Group. Webpage, accessed October 2014. [Online]. Available: http://www.ieee802.org/11/Reports/hew_update.htm
- [17] G. Martorell, F. Riera-Palou, G. Femenias, J. Barcelo, and B. Bellalta, "On the performance evaluation of CSMA/E2CA protocol with open loop ARF-based adaptive modulation and coding," *European Wireless. 18th European Wireless Conference*, pp. 1–8, april 2012.
- [18] E. Yucesan, C. Chen, J. Snowdon, and J. Charnes, "COST: A component-oriented discrete event simulator," in *Winter Simulation Conference*, 2002.
- [19] L. Sanabria-Russo, J. Barcelo, A. Faridi, and B. Bellalta, "WLANs throughput improvement with CSMA/ECA," in *2014 IEEE Conference on Computer Communications Workshops (INFOCOM WKSHPS)*, April 2014, pp. 125–126.
- [20] L. Sanabria-Russo, J. Barcelo, and B. Bellalta, "Prototyping Distributed Collision-Free MAC Protocols for WLANs in Real Hardware," in *Multiple Access Communications*. Springer International Publishing, 2013, pp. 82–87.
- [21] L. Sanabria-Russo. (2013) Report: Prototyping Collision-Free MAC Protocols in Real Hardware. Webpage. [Online]. Available: <http://luissanabria.me/written/beca-test1.pdf>
- [22] L. Sanabria-Russo, F. Gringoli, J. Barcelo and B. Bellalta. (2014) Implementation and Experimental Evaluation of a Collision-Free MAC Protocol for WLANs. Arxiv pre-print. [Online]. Available: <http://arxiv.org/pdf/1410.7924v1>
- [23] F. Gringoli and L. Nava. Open Firmware for WiFi Networks. Webpage. [Online]. Available: <http://www.ing.unibs.it/openfwfw/>
- [24] L. Hanzo, H. Haas, S. Imre, D. O'Brien, M. Rupp, and L. Gyongyosi, "Wireless Myths, Realities, and Futures: From 3G/4G to Optical and Quantum Wireless," *Proceedings of the IEEE*, vol. 100, no. Special Centennial Issue, pp. 1853–1888, 2012.
- [25] J. Lee and J. C. Walrand, "Design and nalysis of an asynchronous zero collision MAC protocol," *arXiv preprint arXiv:0806.3542*, 2008.
- [26] M. Fang, D. Malone, K. Duffy, and D. Leith, "Decentralised learning MACs for collision-free access in WLANs," *Wireless Networks*, pp. 1–16, 2010.
- [27] D. Skordoulis, Q. Ni, H.-H. Chen, A. Stephens, C. Liu, and A. Jamalipour, "IEEE 802.11n MAC Frame Aggregation Mechanisms for Next-Generation High-Throughput WLANs," *Wireless Communications, IEEE*, vol. 15, no. 1, pp. 40–47, February 2008.
- [28] L. Sanabria-Russo. (2014) Github repository: CSMA-ECA-bounds-example. Webpage, accessed October 2014. [Online]. Available: <https://github.com/SanabriaRusso/CSMA-ECA-bounds-example>
- [29] W. Gong and D. Malone, "Addressing Slot Drift in Decentralized Collision Free Access Schemes for WLANs," in *Multiple Access Communications*, ser. Lecture Notes in Computer Science, B. Bellalta, A. Vinel, M. Jonsson, J. Barcelo, R. Maslennikov, P. Chatzimisios, and D. Malone, Eds. Springer Berlin Heidelberg, 2012, vol. 7642, pp. 146–157. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-34976-8_16
- [30] L. Sanabria-Russo, J. Barcelo, and B. Bellalta. (2012) Implementing CSMA/ECA in COST. [Online]. Available: <https://github.com/SanabriaRusso/CSMA-E2CA/tree/journal-branch>
- [31] K. Amiri, Y. Sun, P. Murphy, C. Hunter, J. R. Cavallaro, and A. Sabharwal, "WARP, a Modular Testbed for Configurable Wireless Network Research at Rice," *IEEE SWRIF*, 2007.
- [32] P. Salvador, V. Mancuso, F. Gringoli, and A. Banchs, "VoIPiggy: Analysis and Implementation of a Mechanism to Boost Capacity in IEEE 802.11 WLANs Carrying VoIP Traffic," *IEEE Transactions on Mobile Computing*, vol. 7, no. 13, 2014.
- [33] P. Salvador, L. Cominardi, F. Gringoli, and P. Serrano, "A first implementation and evaluation of the IEEE 802.11aa group addressed transmission service," *Computer Communication Review*, vol. 44, no. 1, 2014.
- [34] D. Berger, F. Gringoli, N. Facchi, I. Martinovic, and J. Schmitt, "Gaining Insight on Friendly Jamming in a Real-world IEEE 802.11 Network," in *Proceedings of the ACM Conference on Security and Privacy in Wireless & Mobile Networks*, 2014.
- [35] A. Tirumala, F. Qin, J. Dugan, J. Ferguson, and K. Gibbs, "Iperf: The TCP/UDP bandwidth measurement tool," <http://dast.nlanr.net/Projects>, 2005.
- [36] J. Barcelo, B. Bellalta, C. Cano, A. Sfairopoulou, M. Oliver, and J. Zuidweg, "Traffic Prioritization for Carrier Sense Multiple Access with Enhanced Collision Avoidance," in *IEEE International Conference on Communications Workshops, 2009. ICC Workshops 2009.*, June 2009, pp. 1–5.